

지능시스템

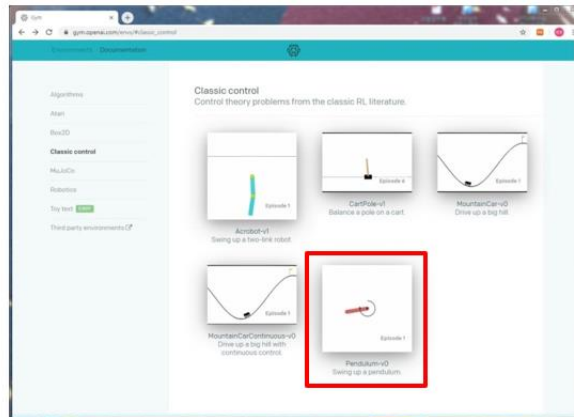
기말대체 과제

2019305059

이현수

학기과제 (DQN 학습 – 도립진자의 제어)

- Cart Pole System(Open AI gym/CartPole-v1)에 대한 DQN 구조의 강화 학습 예제 코드를 참고하여 Open AI gym이 제공하는 Pendulum System을 환경(대상)으로 하여 Pendulum을 수직으로 세우는 제어를 강화학습을 통해 학습시키는 프로그램을 작성하고 결과를 제시하시오. (효과적인 학습을 위한 방안(예를 들어, 보상의 설계)과 제어성능의 우수성을 결과에 포함하여 제시하시오) 6월 14일 월요일까지 Google Classroom에 제출



작업공간(이산, 연속) 및 관측공간(이산, 연속)을 보고 알고리즘을 분류할 수 있다.

1. Q-learning은 개별 작업 공간과 개별 관측 공간을 처리할 수 있다.
2. DQN은 개별 작업 공간과 이산 및 연속 관측 공간을 모두 처리할 수 있다.
3. DDPG는 연속 작업 공간과 이산 및 연속 관측 공간을 모두 처리할 수 있다.

진자 환경과 같이 지속적인 작업 공간이있는 경우 Q-learning 만으로는 연속 작업 공간에서 행동하는 방법을 배울 수 없으므로 정책 그라데이션 모델을 ddpg로 사용해야 한다.

Q-learning에서는 Q-값을 극대화하는 액션을 선택하려고 한다. 이때 단순한 방법으로 가능한 각 작업을 하나씩 살펴보고 Q 값이 가장 높은 것을 선택하는 것이다. 그러나 연속 공간에서는 무한한 양의 가능한 작업이므로 이 방법을 사용할 수 없다.

지속적인 작업을 처리하는 방법에는 여러 가지가 있다. 정책 그라데이션 방법, 소프트 액터 비평가, DDPG 등과 같은 많은 알고리즘이 있다.

이번 과제에서는 DDPG를 사용해 구현했다.

DDPG(심도 결정적 정책 그라데이션)는 연속적인 작업을 학습하기 위한 모델이 없는 오프정책 알고리즘이다.

DPG(결정적 정책 그라데이션)와 **DQN**(Deep Q-Network)의 아이디어를 결합.

DQN 의 익스피리언스 리플레이 및 슬로우 러닝 타겟 네트워크를 사용하며, 연속 작업 공간을 통해 작동할 수 있는 DPG 를 기반.

DDPG 는 2 개의 네트워크를 보유하여 작동한다.

첫 번째 네트워크는 Q-값을 예측하는 법을 학습하고

두 번째 네트워크는 주어진 상태 작업을 선택하는 방법을 학습한다.

<코드>

```
1 import gym
2 import tensorflow as tf
3 from tensorflow.keras import layers
4 import numpy as np
5 import matplotlib.pyplot as plt
6
7 problem = "Pendulum-v0"
8 env = gym.make(problem)
9
10 num_states = env.observation_space.shape[0]
11 print("Size of State Space -> {}".format(num_states))
12 num_actions = env.action_space.shape[0]
13 print("Size of Action Space -> {}".format(num_actions))
14
15 upper_bound = env.action_space.high[0]
16 lower_bound = env.action_space.low[0]
17
18 print("Max Value of Action -> {}".format(upper_bound))
19 print("Min Value of Action -> {}".format(lower_bound))
20
21 class OUActionNoise:
22     def __init__(self, mean, std_deviation, theta=0.15, dt=1e-2, x_initial=None):
23         self.theta = theta
24         self.mean = mean
25         self.std_dev = std_deviation
26         self.dt = dt
27         self.x_initial = x_initial
28         self.reset()
29
30     def __call__(self):
31         x = (
32             self.x_prev
33             + self.theta * (self.mean - self.x_prev) * self.dt
34             + self.std_dev * np.sqrt(self.dt) * np.random.normal(size=self.mean.shape)
35         )
36         self.x_prev = x
37         return x
38
39     def reset(self):
40         if self.x_initial is not None:
41             self.x_prev = self.x_initial
42         else:
43             self.x_prev = np.zeros_like(self.mean)
44
```

```

45 class Buffer:
46     def __init__(self, buffer_capacity=100000, batch_size=64):
47         self.buffer_capacity = buffer_capacity
48         self.batch_size = batch_size
49
50         self.buffer_counter = 0
51
52         self.state_buffer = np.zeros((self.buffer_capacity, num_states))
53         self.action_buffer = np.zeros((self.buffer_capacity, num_actions))
54         self.reward_buffer = np.zeros((self.buffer_capacity, 1))
55         self.next_state_buffer = np.zeros((self.buffer_capacity, num_states))
56
57     def record(self, obs_tuple):
58         index = self.buffer_counter % self.buffer_capacity
59
60         self.state_buffer[index] = obs_tuple[0]
61         self.action_buffer[index] = obs_tuple[1]
62         self.reward_buffer[index] = obs_tuple[2]
63         self.next_state_buffer[index] = obs_tuple[3]
64
65         self.buffer_counter += 1
66
67     @tf.function
68     def update(
69         self, state_batch, action_batch, reward_batch, next_state_batch,
70     ):
71
72         with tf.GradientTape() as tape:
73             target_actions = target_actor(next_state_batch, training=True)
74             y = reward_batch + gamma * target_critic(
75                 [next_state_batch, target_actions], training=True
76             )
77             critic_value = critic_model([state_batch, action_batch], training=True)
78             critic_loss = tf.math.reduce_mean(tf.math.square(y - critic_value))
79
80             critic_grad = tape.gradient(critic_loss, critic_model.trainable_variables)
81             critic_optimizer.apply_gradients(
82                 zip(critic_grad, critic_model.trainable_variables)
83             )
84
85         with tf.GradientTape() as tape:
86             actions = actor_model(state_batch, training=True)
87             critic_value = critic_model([state_batch, actions], training=True)
88
89             actor_loss = -tf.math.reduce_mean(critic_value)
90
91             actor_grad = tape.gradient(actor_loss, actor_model.trainable_variables)
92             actor_optimizer.apply_gradients(
93                 zip(actor_grad, actor_model.trainable_variables)
94             )
95
96

```

```

97     def learn(self):
98         record_range = min(self.buffer_counter, self.buffer_capacity)
99         batch_indices = np.random.choice(record_range, self.batch_size)
100
101         state_batch = tf.convert_to_tensor(self.state_buffer[batch_indices])
102         action_batch = tf.convert_to_tensor(self.action_buffer[batch_indices])
103         reward_batch = tf.convert_to_tensor(self.reward_buffer[batch_indices])
104         reward_batch = tf.cast(reward_batch, dtype=tf.float32)
105         next_state_batch = tf.convert_to_tensor(self.next_state_buffer[batch_indices])
106
107         self.update(state_batch, action_batch, reward_batch, next_state_batch)
108
109
110     @tf.function
111     def update_target(target_weights, weights, tau):
112         for (a, b) in zip(target_weights, weights):
113             a.assign(b * tau + a * (1 - tau))
114
115     def get_actor():
116         last_init = tf.random_uniform_initializer(minval=-0.003, maxval=0.003)
117
118         inputs = layers.Input(shape=(num_states,))
119         out = layers.Dense(256, activation="relu")(inputs)
120         out = layers.Dense(256, activation="relu")(out)
121         outputs = layers.Dense(1, activation="tanh", kernel_initializer=last_init)(out)
122
123         outputs = outputs * upper_bound
124         model = tf.keras.Model(inputs, outputs)
125         return model
126
127
128     def get_critic():
129         state_input = layers.Input(shape=(num_states))
130         state_out = layers.Dense(16, activation="relu")(state_input)
131         state_out = layers.Dense(32, activation="relu")(state_out)
132
133         action_input = layers.Input(shape=(num_actions))
134         action_out = layers.Dense(32, activation="relu")(action_input)
135
136         concat = layers.Concatenate()([state_out, action_out])
137
138         out = layers.Dense(256, activation="relu")(concat)
139         out = layers.Dense(256, activation="relu")(out)
140         outputs = layers.Dense(1)(out)
141
142         model = tf.keras.Model([state_input, action_input], outputs)
143
144         return model
145
146     def policy(state, noise_object):
147         sampled_actions = tf.squeeze(actor_model(state))
148         noise = noise_object()
149         sampled_actions = sampled_actions.numpy() + noise
150
151         legal_action = np.clip(sampled_actions, lower_bound, upper_bound)
152
153         return [np.squeeze(legal_action)]
154

```

```

155 std_dev = 0.2
156 ou_noise = OUActionNoise(mean=np.zeros(1), std_deviation=float(std_dev) * np.ones(1))
157
158 actor_model = get_actor()
159 critic_model = get_critic()
160
161 target_actor = get_actor()
162 target_critic = get_critic()
163
164 target_actor.set_weights(actor_model.get_weights())
165 target_critic.set_weights(critic_model.get_weights())
166
167 critic_lr = 0.002
168 actor_lr = 0.001
169
170 critic_optimizer = tf.keras.optimizers.Adam(critic_lr)
171 actor_optimizer = tf.keras.optimizers.Adam(actor_lr)
172
173 total_episodes = 100
174 gamma = 0.99
175 tau = 0.005
176
177 buffer = Buffer(50000, 64)
178
179 ep_reward_list = []
180 avg_reward_list = []
181
182 for ep in range(total_episodes):
183
184     prev_state = env.reset()
185     episodic_reward = 0
186
187     while True:
188
189         tf_prev_state = tf.expand_dims(tf.convert_to_tensor(prev_state), 0)
190
191         action = policy(tf_prev_state, ou_noise)
192
193         state, reward, done, info = env.step(action)
194
195         buffer.record((prev_state, action, reward, state))
196         episodic_reward += reward
197
198         buffer.learn()
199         update_target(target_actor.variables, actor_model.variables, tau)
200         update_target(target_critic.variables, critic_model.variables, tau)
201
202         if done:
203             break
204
205     prev_state = state
206
207     ep_reward_list.append(episodic_reward)
208
209     avg_reward = np.mean(ep_reward_list[-40:])
210     print("Episode * {} * Avg Reward is ==> {}".format(ep, avg_reward))
211     avg_reward_list.append(avg_reward)
212
213

```

```


214 plt.plot(avg_reward_list)
215 plt.xlabel("Episode")
216 plt.ylabel("Avg. Epsiodic Reward")
217 plt.show()
218
219 actor_model.save_weights("pendulum_actor.h5")
220 critic_model.save_weights("pendulum_critic.h5")
221
222 target_actor.save_weights("pendulum_target_actor.h5")
223 target_critic.save_weights("pendulum_target_critic.h5")
224
225
226 input('press a key to continue for test run')
227 state = env.reset()
228 done = False
229 while not done:
230     tf_prev_state = tf.expand_dims(tf.convert_to_tensor(state), 0)
231     action = policy(tf_prev_state, ou_noise)
232     next_state, reward, done, _ = env.step(action)
233     state = next_state
234     env.render()
235 env.close()
236

```

직접 추가한 코드

학습결과를 render메소드를 통해 비주얼로 표현한다.

Ex)





<실행결과>

```
Run: pendulum
"C:\Program Files\Python38\python.exe" "D:/배달화면/서경대공부(5학기)/5학기/지능시스템/과제(기밀대제)/pendulum.py
2021-06-08 09:41:06.314158: W tensorflow/stream_executor/platform/default/dso_loader.cc:60] Could not load dynamic library 'cudart64_110.dll'; dlderror: cudart64_110.dll not found
2021-06-08 09:41:06.314298: I tensorflow/stream_executor/cuda/cudart_stub.cc:29] Ignore above cudart dlderror if you do not have a GPU set up on your machine.
Size of State Space -> 3
Size of Action Space -> 1
Max Value of Action -> 2.0
Min Value of Action -> -2.0
2021-06-08 09:41:09.063978: I tensorflow/compiler/jit/xla_cpu_device.cc:41] Not creating XLA devices, tf_xla_enable_xla_devices not set
2021-06-08 09:41:09.066402: W tensorflow/stream_executor/platform/default/dso_loader.cc:60] Could not load dynamic library 'nvcuda.dll'; dlderror: nvcuda.dll not found
2021-06-08 09:41:09.066572: W tensorflow/stream_executor/cuda/cuda_driver.cc:326] failed call to cuInit: UNKNOWN ERROR (303)
2021-06-08 09:41:09.079851: I tensorflow/stream_executor/cuda/cuda_diagnostics.cc:169] retrieving CUDA diagnostic information for host: DESKTOP-OIMG8KE
2021-06-08 09:41:09.080181: I tensorflow/stream_executor/cuda/cuda_diagnostics.cc:176] hostname: DESKTOP-OIMG8KE
2021-06-08 09:41:09.080845: I tensorflow/core/platform/cpu_feature_guard.cc:142] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instruction
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
2021-06-08 09:41:09.081647: I tensorflow/compiler/jit/xla_gpu_device.cc:99] Not creating XLA devices, tf_xla_enable_xla_devices not set
2021-06-08 09:41:09.942236: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:116] None of the MLIR optimization passes are enabled (registered 2)
Episode * 0 * Avg Reward is ==> -1266.904386491743
Episode * 1 * Avg Reward is ==> -1160.1012719445998
Episode * 2 * Avg Reward is ==> -1262.942140479222
Episode * 3 * Avg Reward is ==> -1401.4365392073937
Episode * 4 * Avg Reward is ==> -1458.4344015262475
Episode * 5 * Avg Reward is ==> -1470.4876812372083
Episode * 6 * Avg Reward is ==> -1469.1989289819337
Episode * 7 * Avg Reward is ==> -1472.7778237011285
Episode * 8 * Avg Reward is ==> -1408.9014248705666
Episode * 9 * Avg Reward is ==> -1374.2343110039183
Episode * 10 * Avg Reward is ==> -1309.071241016784
```

```
Episode * 0 * Avg Reward is ==> -981.2168329701693
Episode * 1 * Avg Reward is ==> -974.5008333503071
Episode * 2 * Avg Reward is ==> -1047.048232919383
Episode * 3 * Avg Reward is ==> -1097.1185316541814
Episode * 4 * Avg Reward is ==> -1147.2553260405266
Episode * 5 * Avg Reward is ==> -1186.6385431023948
Episode * 6 * Avg Reward is ==> -1234.2580864198005
Episode * 7 * Avg Reward is ==> -1243.0193229597623
Episode * 8 * Avg Reward is ==> -1268.8086857653263
Episode * 9 * Avg Reward is ==> -1254.7276127314738
Episode * 10 * Avg Reward is ==> -1238.7468729280206
Episode * 11 * Avg Reward is ==> -1237.682666878154
Episode * 12 * Avg Reward is ==> -1192.6486879129957
Episode * 13 * Avg Reward is ==> -1117.255529263095
Episode * 14 * Avg Reward is ==> -1051.5623556277592
Episode * 15 * Avg Reward is ==> -1002.1831328285903
Episode * 16 * Avg Reward is ==> -951.0421874306827
Episode * 17 * Avg Reward is ==> -912.5647587025315
Episode * 18 * Avg Reward is ==> -877.6033536272961
Episode * 19 * Avg Reward is ==> -846.6414423935597
Episode * 20 * Avg Reward is ==> -824.9676317441287
Episode * 21 * Avg Reward is ==> -793.2398441250187
Episode * 22 * Avg Reward is ==> -769.6546150419948
Episode * 23 * Avg Reward is ==> -742.4887919541511
Episode * 24 * Avg Reward is ==> -724.3730749541766
Episode * 25 * Avg Reward is ==> -747.5590212096414
Episode * 26 * Avg Reward is ==> -724.2935983092336
Episode * 27 * Avg Reward is ==> -751.283781700285
Episode * 28 * Avg Reward is ==> -729.6051407181939
Episode * 29 * Avg Reward is ==> -713.381735752476
```

(중간 생략)

```
Episode * 90 * Avg Reward is ==> -214.68677771191068
Episode * 91 * Avg Reward is ==> -209.46807976530923
Episode * 92 * Avg Reward is ==> -206.61452482715077
Episode * 93 * Avg Reward is ==> -203.57159869114548
Episode * 94 * Avg Reward is ==> -203.60756337403458
Episode * 95 * Avg Reward is ==> -203.6327072714442
Episode * 96 * Avg Reward is ==> -206.70653307738075
Episode * 97 * Avg Reward is ==> -203.82494869795292
Episode * 98 * Avg Reward is ==> -203.95345773492153
Episode * 99 * Avg Reward is ==> -201.46951187358368
```



결과를 보면 학습이 정상적으로 되어 목표하는 결과가 나왔다.